

Document Classification using NLP

TF-IDF Representation + Machine Learning

Course: Natural Language Processing | Assignment 1

Dataset:	trainset__3_.txt (Reuters News Corpus, Binary Classification)
Total Samples:	200 documents (+1: 102 Agriculture -1: 98 Non-Agriculture)
Features:	Custom TF-IDF (unigram+bigram, sublinear_tf, min_df=2, max_features=5000) → 3,126 features
Algorithm:	Logistic Regression / Linear SVM / Multinomial Naive Bayes
Evaluation:	5-Fold Stratified Cross-Validation + 80/20 Hold-out Test

1. Introduction

This report presents a complete document classification pipeline applied to the Reuters News Corpus. The task is binary classification: given a news article, determine whether it belongs to the **Agriculture (+1)** category (covering grain, sugar, cocoa, coffee, oilseeds, crop reports, etc.) or the **Non-Agriculture (-1)** category (covering metals, mining, steel, industrial commodities, and financial news).

The dataset contains 200 labeled documents with a near-balanced class distribution (51% positive, 49% negative), making it suitable for standard accuracy and F1 evaluation without class-weighting adjustments.

Two anomalous documents (advance-fee fraud / spam emails) appear in the dataset with a -1 label — these are handled naturally by the classifier since their vocabulary is entirely different from commodity news.

2. Dataset Analysis

2.1 Class Distribution

Class	Label	Count	Percentage
Agriculture	+1	102	51.0%
Non-Agriculture	-1	98	49.0%
Total	—	200	100%

2.2 Sample Documents

Agriculture (+1): "ROTTERDAM GRAIN HANDLER SAYS PORT BALANCE ROSE — Graan Elevator Mij, GEM, said its balance in port of grains, oilseeds and derivatives rose to 146,000 tonnes..."

Non-Agriculture (-1): "NIPPON MINING LOWERS COPPER PRICE — Nippon Mining Co Ltd said it lowered its selling price for electrolytic copper by 10,000 yen per tonne to 260,000..."

3. Custom Text Representation

We designed a custom TF-IDF representation that goes beyond basic bag-of-words by incorporating several evidence-based design choices suited to news classification:

Design Choice	Value	Justification
N-gram range	Unigrams + Bigrams (1,2)	Captures commodity phrases: "grain exports", "copper prices", "sugar production"
TF scaling	Sublinear (1 + log tf)	Prevents dominance of high-frequency agricultural jargon
Vocabulary cap	5,000 features	Controls model complexity; prevents overfitting on rare terms
Min. document freq.	2	Filters one-off noise: typos, company codes, date strings
Stop words	English stop words removed	Removes uninformative function words (the, is, at, etc.)
IDF weighting	Smooth IDF (default)	Down-weights terms common to both classes

3.1 Feature Statistics

The resulting feature matrix has **3,126 features** across **200 documents**. Representative discriminative bigrams identified include:

- **Agriculture (+1):** grain exports, wheat barley, corn soybean, coffee quota, sugar production, crop report, agriculture department
- **Non-Agriculture (-1):** copper price, gold mine, zinc tonnes, steel corp, lead production, mining inc

3.2 Preprocessing Pipeline

Each document undergoes the following preprocessing before vectorization:

1. Lowercase conversion
2. HTML entity removal (e.g., <COMPANY> → removed)
3. Non-alphabetic removal (numbers, punctuation stripped)
4. Whitespace normalization
5. English stop word removal (handled by TfidfVectorizer)
6. TF-IDF with sublinear scaling applied

4. Classifier Selection and Methodology

Three classifiers were selected to cover a range of algorithmic approaches, all well-suited to high-dimensional sparse text data:

Model	Hyperparameters	Rationale
Logistic Regression	C=1.0, max_iter=1000, L2 penalty	Highly interpretable baseline; coefficient magnitudes directly indicate feature importance
Linear SVM	C=1.0, kernel=linear, probability=False	State-of-the-art for text; maximises margin in high-dim sparse feature space
Multinomial Naive Bayes	alpha=0.1 (Laplace smoothing)	Probabilistic baseline; fast; works well with TF-IDF counts

4.1 Evaluation Protocol

A two-stage evaluation was used to ensure robust performance estimation:

Stage 1 — 5-Fold Stratified Cross-Validation: All 200 documents are split into 5 folds with class-ratio preservation. Each fold serves once as the test set. Mean $\pm$ std accuracy and F1 are reported.

Stage 2 — 80/20 Hold-out Evaluation: A fixed 80/20 stratified split provides precision, recall, F1 per class, confusion matrix, and ROC-AUC for detailed error analysis.

5. Results

5.1 Cross-Validation Results

Model	CV Accuracy	CV F1 Score	F1 Std Dev
Logistic Regression	0.9550	0.9571	$\pm$ 0.0225
SVM (Linear)	0.9850	0.9856	$\pm$ 0.0191
Naive Bayes	0.9600	0.9623	$\pm$ 0.0232

** Highlighted row = best model (SVM Linear)*

5.2 Hold-out Test Set Results (80/20 Split)

The best model (SVM Linear) was evaluated on the held-out test set (40 documents):

Metric	Agriculture (+1)	Non-Agriculture (-1)	Overall
Precision	0.947	0.905	0.925
Recall	0.900	0.950	0.923
F1-Score	0.923	0.927	0.923
ROC-AUC	—	—	0.9875
Accuracy	—	—	0.9250

5.3 Confusion Matrix

	Predicted: Non-Agric (-1)	Predicted: Agriculture (+1)
True: Non-Agric (-1)	19	1
True: Agriculture (+1)	2	18

TP=18 (correct Agriculture) | TN=19 (correct Non-Agriculture) | FP=1 (false alarm) | FN=2 (missed Agriculture)

6. Visualisations




Figure 1: (Left) Model comparison across 5-fold CV. (Centre) Confusion matrix on test set. (Right) ROC curve with AUC score.

7. Error Analysis

With only 3 misclassifications out of 40 test documents, the error rate is very low. Analysis of the errors reveals interesting boundary cases:

False Positives (1): Non-agriculture documents predicted as agriculture. These likely involve articles about agricultural-related finance or trade policy that share vocabulary with agricultural news.

False Negatives (2): Agriculture documents predicted as non-agriculture. These tend to be short or title-only records (NO TEXT), making classification ambiguous based solely on the headline.

The two spam/fraud emails in the dataset (labelled -1) were correctly classified as non-agriculture since their vocabulary ("bank account", "transfer", "USD million") is absent from agricultural documents.

8. Discussion

Why TF-IDF with Bigrams Works: The Reuters news corpus has highly topic-specific vocabulary. Terms like "wheat", "soybean", "cocoa" are almost exclusively agricultural, while "copper", "zinc", "mining" characterise the negative class. TF-IDF amplifies these discriminative terms and down-weights common news-writing vocabulary.

Why SVM Outperforms: Linear SVMs are theoretically well-suited to text because: (1) they find the maximum-margin hyperplane in the high-dimensional TF-IDF space; (2) they are robust to irrelevant features; (3) they handle sparse input efficiently. With 3,126 features and only 200 samples, the margin-maximisation objective prevents overfitting.

Bigrams vs. Unigrams: Preliminary experiments showed bigrams added 0.5–1.5% F1 improvement because they capture compound commodity names ("grain exports", "coffee quota") that are uniquely discriminative.

Limitations: (1) Small dataset (200 documents) — results may vary with new data. (2) No lemmatization applied — "farms" and "farming" treated as different features. (3) No hyperparameter tuning (grid search) performed due to dataset size.

9. Conclusion

Aspect	Outcome
Best Model	Linear SVM (TF-IDF features)
CV Accuracy	98.50% $\pm$ 2.00%
CV F1 Score	98.56% $\pm$ 1.91%
Test Accuracy	92.50% (on held-out 40 docs)
ROC-AUC	0.9875 (near-perfect class separation)
Feature Space	3,126 TF-IDF features (unigram + bigram)
Key Finding	Bigrams + sublinear TF are the most important representation choices

The classification pipeline demonstrates that a carefully designed TF-IDF representation combined with a Linear SVM achieves near-perfect performance on this Reuters corpus binary classification task. The success highlights the importance of domain-appropriate feature engineering — particularly the use of bigrams and sublinear TF scaling — over raw unigram bag-of-words approaches. Future work could explore deep learning approaches (e.g., fine-tuned BERT) and extend to multi-class topic classification across all Reuters categories.